The LLVM compiler framework

Introduction

Daniele Cattaneo

Politecnico di Milano

2020-05-20

These slides were originally written by Michele Scandale, Ettore Speziale and Stefano Cherubin for the "Code Transformation and Optimization" course.

Contents

Compilers and compilers

Not all compilers are the same... The traditional distinction was:

Toy Compiler	Production-Quality Compiler
▶ small codebase	▶ huge codebase
▶ easy to modify	▶ hard to modify
▶ limited capabilities	▶ produces high-quality code

Working with a production-quality compiler is *initially* **hard**...

...however it provides a huge set of tools that toy compilers **miss**!

LLVM: Low Level Virtual Machine

Initially started as a small research project at Urbana-Champaign.

Now it has grown to a huge size...

► Key technology in the **industry**: AMD, Apple, Google, Intel, NVIDIA...

► Still intensively used in **research** about compilers

GCC vs LLVM

LLVM [1] is Open Source

If you are familiar with Linux you might have used **GCC** [2]...

GCC is older than LLVM

$\Rightarrow$ GCC produces better code

$\Rightarrow$ LLVM is generally faster

$\Rightarrow$ LLVM is more modular and *clean*

Contents

# Compiler pipeline

Typically a compiler is a **pipeline**.

Advantages of the pipeline model:

► **simplicity** – read something, produce something
► **locality** – no superfluous state

Complexity lies on **chaining** together stages.

Compiler pipeline

High-level pipeline structure of a compiler:

There are three main components:

Front-end **translate** a source file into the intermediate representation

Middle-end **analyze** the intermediate representation, **optimize** it

Back-end **generate** target machine assembly from the intermediate representation

The LLVM compiler pipeline

We will consider only the *middle-end*.
Same concepts are valid also for {front,back}-end.

IR (a.k.a. Intermediate Representation) the **language** used in the middle-end

Pass a **pipeline stage**
a Pass may have **dependencies** on other Passes.

Pass Manager component that **schedules** passes according to their **dependencies** and **executes** them
(builds the pipeline)

First insights

A compiler is **complex**:

- ▶ passes are the **elementary unit of work**
- ▶ Pass Manager must be **advised** about pass chaining
- ▶ pipeline shapes are **not fixed** – they can change from one compiler execution to another
  e.g. optimized/not optimized builds, compiler options, ...

A word of warning!

Compilers must be **conservative**:

⇓

All passes **must preserve the program semantics**

⇓

Compiler passes must be designed **very carefully**!

Contents

Classical Algorithm Design

In algorithm design, a good approach is the following:

1. study the problem
2. make some example
3. identify the **common case**
4. derive the algorithm for the common case
5. add handling for **corner cases**
6. improve performancing **optimizing the common case**

Weakness of the approach:

▶ **corner cases** – a *correct* algorithm **must** consider *all the corner cases*!

Compiler Algorithm Design

Corner cases are difficult to handle, but they cannot be ignored

Compiler algorithms must be **proven** to preserve
program semantic **at all times**

As an aid, a *standard methodology* is employed.

Compiler algorithms are built combining **three** kinds of passes:

Analysis, Optimization, Normalization

Compiler Algorithm Design

Corner cases are difficult to handle, but they cannot be ignored

Compiler algorithms must be **proven** to preserve
program semantic **at all times**

As an aid, a *standard methodology* is employed.

Compiler algorithms are built combining **three** kinds of passes:

Analysis, Optimization, Normalization

We now consider a simple example: *loop hoisting*.

Loop Hoisting

It is a transformation that:

- ► looks for statements (inside a loop) not depending on the loop state
- ► move them outside the loop body

Loop Hoisting – Before	Loop Hoisting – After

```
do {
  a += i;
  b = c;
  i++;
} while (i < k);
```

```
b = c;
do {
  a += i;
  i++;
} while (i < k);
```

Loop Hoisting

The general idea:

- ► move "good" statement outside of the loop

This **pass** modifies the code, thus it is an **optimization pass**.
It needs to know:

- ► which pieces of code are loops
- ► which statements are "good" statements

This information is computed by the the **analysis** passes:

- ► detecting loops in the program
- ► detecting loop-independent statements

The loop hoisting pass declares which analyses it needs:

- ► pipeline automatically built: **analysis → optimization**

Loop Hoisting

The **proof** is trivial:

- ▶ the transformation shall preserve program semantics
- ▶ the analyses shall be correct

Analysis passes are usually built starting from other analyses already implemented inside the compiler, or are already present in LLVM

- ▶ often, no proof is necessary for the analyses

However...

You also have to prove that the combination of analysis + transformation is correct!

"Beware of bugs in the above code;
I have only proved it correct, not tried it."
— Knuth

Loop Hoisting

We have spoken about loops, but which kind of loop?

do-while? **while**? **for**?

Almost all loops are different forms of the **same exact thing**
⇓
We can convert a lot of loops to a loop of another kind!

To account for the various kinds of loops, we choose a **normal** kind of
loop, and then we write a **normalization** pass.

Usually, **do-while** loops are chosen to be the *normal* loops.
Sometimes, normalization is also called **canonicalization**.

The more loops we recognize, the higher the potential
optimization impact!

Compiler Algorithm Design

You have to:

1. analyze the problem
2. make some examples
3. detect the common case
4. determine the **input conditions**
5. determine which **analyses** you need
6. design the **optimization** pass
7. proof its **correctness**
8. improve algorithm perfomance on the common case
9. improve the effectiveness of the algorithm by adding **normalization passes**

Compiler Algorithm Design

Something is missing...

Corner Cases!

Why?

1. It makes no sense to optimize code that is seldom executed
2. Your optimization will be based on **properties of the code that are true only in the common case you are considering**
   - ► If the code does not fit the common case, it shall stay as-is
   - ► Otherwise you **risk breaking program semantics**!

Contents

Using LLVM

LLVM is a **compiler construction framework**
It operates on the **LLVM-IR** language.

⇓

Using LLVM *by itself* does not make much sense!
Writing LLVM-IR by hand is unfeasible.

Terminology

LLVM-IR comes in 3 different flavours:

assembly on-disk human-readable format
(file extension: `.ll`)

bitcode on-disk machine-oriented binary format
(file extension: `.bc`)

in-memory in-memory binary format
(used during compilation process)

All formats have the same expressiveness!

Frontends and Drivers

The LLVM-IR input to LLVM is provided by **frontends**.

Example

Clang[3] is the frontend for the C language family

The **compiler driver** is the program that:
- ▶ Provides the interface to the user
- ▶ Performs setup of the front end and LLVM itself.

Example

The driver of *Clang* is the `clang` executable (compatible with GCC)

Frontends and Drivers

We can generate LLVM-IR assembly using the `clang` driver:

```
clang -emit-llvm -S -o out.ll in.c
```

If you want to generate bitcode instead:
```
clang -emit-llvm -o out.bc in.c
```

The compiler driver can also generate native code starting from
LLVM-IR assembly

(Like compiling an assembly file with GCC)

Tools

Run one or more passes on the LLVM-IR on-demand by using `opt`:

- ▶ Syntax is like `clang` (supports even -O1, -O2...)
- ▶ One command line argument per pass to run
- ▶ Order of execution is the same as the argument order
  - ▶ Different order, different results! (**phase/stage ordering**)

Some useful passes for debugging (they do not transform anything):

print CFG `opt -view-cfg input.ll`

print dominator tree `opt -view-dom input.ll`

print current IR `opt -print-module input.ll`

Example

- ▶ Run *mem2reg*, then view the CFG:
  - ▶ `opt -mem2reg -view-cfg input.ll`

# Pass Hierarchy

LLVM provides a lot of passes...

▶ Try `opt -help`!

For performance reasons there are different kind of passes:

LLVM Passes

```
                              Pass
                               △
                               │
   ┌─────────────┬─────────────┼─────────────┬─────────────┐
CallGraphSCCPass ModulePass FunctionPass   LoopPass    RegionPass
                    △
                    │
               ImmutablePass
```

LLVM Passes

Each kind of pass visits particular elements of a module:

ImmutablePass compiler configuration – never run

CallGraphSCCPass post-order visit of CallGraph SCCs

ModulePass visit the whole module

FunctionPass visit functions

LoopPass post-order visit of loop nests

RegionPass visit a custom-defined region of code

Specializations come with restrictions:

- ► e.g. a **FunctionPass** cannot add or delete functions
- ► refer to "Writing a LLVM Pass" [4] for documentation on features and limitations of each kind of pass

Normalization passes in LLVM

LLVM provides several **normalization/canonicalization** passes:

- ► Variable-to-register promotion (`Mem2Reg`)
- ► Loop canonicalization (`LoopSimplify`)
- ► CFG canonicalization & simplification (`SimplifyCFG`)
- ► …

They are useful to make your life easier!

Contents

LLVM IR

The LLVM IR [5] language is RISC-based:
- ▶ instructions operates on **variables**
    - ▶ aka virtual registers, temporary values
- ▶ only **load** and **store** access memory
- ▶ **alloca** used to reserve memory on the stack

There are also a few **high level instructions**:
- ▶ function call – **call**
- ▶ pointer arithmetics – **getelementptr**
- ▶ …

LLVM IR

LLVM IR is **strongly typed**:

 ► e.g. you cannot assign a floating point value to an integer
   variable without an explicit cast

Almost everything is **typed**:

   functions @fact   → i32 (i32)
 statements %3 = **icmp eq** i32 %2, 0   → i1

A variable can be:

      global @var = **common global** i32 0, **align** 4
  argument **define** i32 @fact(i32 %n)
       local %2 = **load** i32, i32* %1, **align** 4

Local variables are defined by statements

LLVM IR

```
define i32 @fact(i32 %n) {
entry:
  %retval = alloca i32, align 4
  %n.addr = alloca i32, align 4
  store i32 %n, i32* %n.addr, align 4
  %0 = load i32, i32* %n.addr, align 4
  %cmp = icmp eq i32 %0, 0
  br i1 %cmp, label %if.then, label %if.end

if.then:
  store i32 1, i32* %retval, align 4
  br label %return

if.end:
  %1 = load i32, i32* %n.addr, align 4
  %2 = load i32, i32* %n.addr, align 4
  %sub = sub nsw i32 %2, 1
  %call = call i32 @fact(i32 %sub)
  %mul = mul nsw i32 %1, %call
  store i32 %mul, i32* %retval, align 4
  br label %return

return:
  %3 = load i32, i32* %retval, align 4
  ret i32 %3
}
```

LLVM IR Language

LLVM IR is SSA-based:

► every variable is **statically assigned** exactly **once**

Statically means that:

► inside each function...

► ...for each variable %foo...

► ...there is **only one** statement in the form %foo = ...

Static (compile time) $\neq$ **dynamic** (runtime)

► Single *Dynamic* Assignment:
   in the execution trace there is only one assignment to a variable x

► Single *Static* Assignment:
   in the code listing there is only one assignment to a variable x

    ► Assignments **can** be performed multiple times (in a loop for example)

Static Single Assignment

Scalar SAXPY

```
float saxpy(float a, float x, float y) {
    return a * x + y;
}
```

Scalar LLVM SAXPY

```
define float @saxpy(float %a, float %x, float %y) {
  %1 = fmul float %a, %x
  %2 = fadd float %1, %y
  ret float %2
}
```

Temporary %1 not reused! %2 is used for the second assignment!

Static Single Assignment

Array SAXPY

```
void saxpy(float a, float x[4], float y[4], float z[4]) {
  for(unsigned i = 0; i < 4; ++i)
    z[i] = a * x[i] + y[i];
}
```

Array LLVM SAXPY

```
for.cond:
  %i.0 = phi i32 [ 0, %entry ], [ %inc, %for.inc ]
  %cmp = icmp ult i32 %i.0, 4
  br i1 %cmp, label %for.body, label %for.end
  [...]
for.inc:
  %inc = add i32 %i.0, 1
  br label %for.cond
```

One assignment for loop counter %i.0

Static Single Assignment

Max

```
float max(float a, float b) {
  return a > b ? a : b;
}
```

LLVM Max – WRONG

```
  %1 = fcmp ogt float %a, %b
  br i1 %1, label %if.then, label %if.else
if.then:
  %2 = %a
  br label %if.end
if.else:
  %2 = %b
  br label %if.end
if.end:
  ret float %2
```

Why is it **wrong**?

Static Single Assignment

The %2 variable must be statically assigned once!

LLVM Max

```
  %1 = fcmp ogt float %a, %b
  br i1 %1, label %if.then, label %if.end
if.then:
  br label %if.end
if.else:
  br label %if.end
if.end:
  %2 = phi float [ %a, %if.then ], [ %b, %if.else ]
  ret float %2
```

The **phi** instruction is a *conditional move*:

► it takes $(\text{variable}_i, \text{label}_i)$ pairs
► if coming from predecessor identified by label_i, its value is variable_i

Static Single Assignment

Each SSA variable is assigned only once:

► variable **definition**

Each SSA variable can be referenced by multiple instructions:

► variable **uses**

Algorithms and technical language abuse of these terms!

Let %foo be a variable. If the definition of %foo does not have side-effects nor uses, the aforementioned %foo variable can be erased from the CFG without altering program semantics.

Static Single Assignment

Old compilers are not SSA-based:
- ► converting non-SSA input into SSA form is expensive
- ► cost must be amortized

New compilers are SSA-based:
- ► SSA easier to work with
- ► SSA-based analysis/optimizations are faster

Contents

Conclusions

LLVM is a **production-quality** compiler framework:

$\Rightarrow$ impossible knowing all details

But:

- ▶ it is well organized
- ▶ if you know compiler theory, it is relatively easy to find what you need inside the source code

Remember it's written in C++!

- ▶ To hack around LLVM you need at least basic C++ skills
- ▶ C++ $\neq$ C

Contents

Bibliography I

University of Illinois at Urbana-Champaign.
Low Level Virtual Machine.
http://www.llvm.org.

GNU.
GNU Compiler Collection.
http://gcc.gnu.org.

University of Illinois at Urbana-Champaign.
Clang: a C language family frontend for LLVM.
http://clang.llvm.org.

Chris Lattner and Jim Laskey.
Writing an LLVM Pass.
http://llvm.org/docs/WritingAnLLVMPass.html.

Bibliography II

Chris Lattner and Vikram Adve.
LLVM Language Reference Manual.
http://llvm.org/docs/LangRef.html.

Linus Torvalds.
Re: SCO: "thread creation is about a thousand times faster than onnative.
https://lkml.org/lkml/2000/8/25/132.

Bruce Eckel.
Thinking in C++ – Volume One: Introduction to Standard C++.
http://mindview.net/Books/TICPP/ThinkingInCPP2e.html.

Bruce Eckel and Chuck Allison.
Thinking in C++ – Volume Two: Practical Programming.
http://mindview.net/Books/TICPP/ThinkingInCPP2e.html.

Bibliography III

📄 AMD.
Open64.
http://developer.amd.com/tools-and-sdks/cpu-development/x86-open64-compiler-suite.

📄 John T. Criswell, Daniel Dunbar, Reid Spencer, and Tanya Lattner.
LLVM Testing Infrastructure Guide.
http://llvm.org/docs/TestingGuide.html.

📄 LLVM Community.
LLVM Coding Standards.
http://llvm.org/docs/CodingStandards.html.

📄 LLVM Community.
LLVM Passes.
http://llvm.org/docs/Passes.html.

Bibliography IV

📄 LLVM Community.
Autovectorization in LLVM.
http://llvm.org/docs/Vectorizers.html.

📄 LLVM Community.
LLVM Programmer's Manual.
http://llvm.org/docs/ProgrammersManual.html.

📄 Ettore Speziale.
Compiler Optimization and Transformation Passes.
https://github.com/speziale-ettore/COTPasses.

📄 Scott Chacon.
Pro Git.
http://git-scm.com/book.