

# The LLVM compiler framework

## Introduction

Daniele Cattaneo

Politecnico di Milano

2021-03-16

*These slides are based on material originally written by Michele Scandale, Ettore Speziale and Stefano Cherubin for the “Code Transformation and Optimization” course.*

# Contents

**Introduction**

Compiler organization

Algorithm design

Inside LLVM

Conclusions

Bibliography

# Compilers and compilers

You might have already have had experience working in a **toy compiler...**

## Toy Compiler

- ▶ small codebase
- ▶ easy to modify
- ▶ limited capabilities

## Production-Quality Compiler

- ▶ huge codebase
- ▶ hard to modify
- ▶ produces high-quality code

*Initially*, working with a production-quality compiler might seem **hard...**  
...however it provides a huge set of tools that toy compilers **miss!**

# Why LLVM

Why is this course focused on LLVM?

- ▶ Key technology in the **industry**
  - ▶ AMD, Apple, Google, Intel, NVIDIA...
- ▶ Biggest platform for **research** about compilers
- ▶ **Modular** and **hackable**

Initially started as a small research project at Urbana-Champaign.

Now it has grown to a huge size...

# GCC vs LLVM

**LLVM** [1] is Open Source

If you are familiar with Linux you might have used **GCC** [2]...

GCC is older than LLVM

- ⇒ GCC produces better code
- ⇒ LLVM is generally faster
- ⇒ LLVM is more modular and *clean*

# Contents

Introduction

**Compiler organization**

Algorithm design

Inside LLVM

Conclusions

Bibliography

# Compiler pipeline

Typically a compiler is a **pipeline**.

Advantages of the pipeline model:

- ▶ **simplicity** – read something, produce something
- ▶ **locality** – no superfluous state

Complexity lies on **chaining** together stages.

# Frontends and Drivers

The external interface to LLVM is provided by the **compiler driver**.

The **compiler driver** is the program that:

- ▶ Provides the interface to the user
- ▶ Performs setup of the frontend and LLVM itself.

The driver invokes the **first stage of the pipeline**, the **frontend**.

## Example

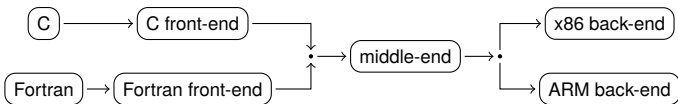
**Clang**[3] is the frontend for the C language family.

The driver of *Clang* is the `clang` executable (compatible with GCC).



# Compiler pipeline

High-level pipeline structure of a compiler:



There are three main components:

**Front-end translates** a source file into the intermediate representation

**Middle-end analyzes** the intermediate representation, **optimizes** it

**Back-end generates** target machine assembly from the intermediate representation

# The LLVM compiler pipeline

We will focus on the *middle-end*.

Same concepts are valid also for {front,back}-end.

- ▶ The *front-end* is completely language-specific
  - ▶ There are no special facilities in LLVM for parsing and AST generation
  - ▶ This is slowly changing: **MLIR**
- ▶ The *back-end* uses its own machine-specific IR
  - ▶ As opposed to the frontend, in LLVM there are generalized facilities for implementing a backend
  - ▶ Complex topic, not enough time...
  - ▶ Most optimizations happen in the middle-end anyway

# The LLVM compiler pipeline

The lingua franca of LLVM is its **Intermediate Representation** called  
LLVM-IR

The LLVM-IR is:

- ▶ **Produced** by the **front-end**
- ▶ **Transformed and optimized** by the **middle-end**
- ▶ **Consumed** by the **back-end**

Understanding LLVM-IR is the key to hacking within LLVM.

# Remember...

LLVM is a **compiler construction framework**  
It operates on the **LLVM-IR** language.



Using LLVM *by itself* does not make much sense!  
Writing LLVM-IR by hand is unfeasible.

# The Middle End

The middle-end is not a monolithic block,  
but it is a pipeline in and of itself.

**LLVM-IR** the **language** used in the middle-end

**Pass** a **pipeline stage**

a Pass may have **dependencies** on other Passes.

**Pass Manager** component that **schedules** passes according to their **dependencies** and **executes** them  
*(builds the pipeline)*

Our focus: **writing a pass**

# First insights

A compiler is **complex**:

- ▶ passes are the **elementary unit of work**
- ▶ Pass Manager must be **advised** about pass chaining
- ▶ pipeline shapes are **not fixed** – they can change from one compiler execution to another  
e.g. optimized/not optimized builds, compiler options, ...

# A word of warning!

Compilers must be **conservative**:



All passes **must preserve the program semantics**



Compiler passes must be designed **very carefully!**

# Contents

Introduction

Compiler organization

**Algorithm design**

Inside LLVM

Conclusions

Bibliography



# Classical Algorithm Design

In algorithm design, a good approach is the following:

1. study the problem
2. make some example
3. identify the **common case**
4. derive the algorithm for the common case
5. add handling for **corner cases**
6. improve performance by **optimizing the common case**

Weakness of the approach:

- ▶ **corner cases:**  
a *correct* algorithm **must** consider *all the corner cases*!

# Compiler Algorithm Design

Corner cases are difficult to handle, but they cannot be ignored

Compiler algorithms must be **proven** to preserve  
program semantic **at all times**

As an aid, a *standard methodology* is employed.

Compiler algorithms are built combining **three** kinds of passes:

Analysis, Optimization, Normalization

# Compiler Algorithm Design

Corner cases are difficult to handle, but they cannot be ignored

Compiler algorithms must be **proven** to preserve  
program semantic **at all times**

As an aid, a *standard methodology* is employed.

Compiler algorithms are built combining **three** kinds of passes:

Analysis, Optimization, Normalization

We now consider a simple example: *loop hoisting*.

# Loop Hoisting

It is a transformation that:

- ▶ looks for statements (inside a loop) not depending on the loop state
- ▶ move them outside the loop body

## Loop Hoisting – Before

```
do {  
    a += i;  
    b = c;  
    i++;  
} while (i < k);
```

## Loop Hoisting – After

```
b = c;  
do {  
    a += i;  
    i++;  
} while (i < k);
```

# Loop Hoisting

## The general idea:

- ▶ move “good” statement outside of the loop

This **pass** modifies the code, thus it is an **optimization pass**.  
It needs to know:

- ▶ which pieces of code are loops
- ▶ which statements are “good” statements

This information is computed by the the **analysis** passes:

- ▶ detecting loops in the program
- ▶ detecting loop-independent statements

The loop hoisting pass declares which analyses it needs:

- ▶ pipeline automatically built: **analysis** → **optimization**

# Loop Hoisting

The **proof** is trivial:

- ▶ the transformation shall preserve program semantics
- ▶ the analyses shall be correct

Analysis passes are usually built starting from other analyses already implemented inside the compiler, or are already present in LLVM

- ▶ often, no proof is necessary for the analyses

## However...

You also have to prove that the combination of analysis + transformation is correct!

"Beware of bugs in the above code;  
I have only proved it correct, not tried it."

— Knuth

# The importance of normalization

We have spoken about loops, but which kind of loop?

**do-while? while? for?**

Almost all loops are different forms of the **same exact thing**



We can convert a lot of loops to a loop of another kind!

To account for the various kinds of loops, we choose a **normal** kind of loop, and then we write a **normalization** pass.

Usually, **do-while** loops are chosen to be the *normal* loops.

Sometimes, normalization is also called **canonicalization**.

The more loops we recognize, the higher the potential **optimization impact!**

# A Methodology

You have to:

1. analyze the problem
2. make some examples
3. detect the common case
4. determine the **input conditions**
5. determine which **analyses** you need
6. design the **optimization** pass
7. proof its **correctness**
8. improve algorithm performance on the common case
9. improve the effectiveness of the algorithm by adding **normalization passes**



# Ignore corner cases!

Something is missing...

## Corner Cases!

Why?

1. It makes no sense to optimize code that is seldom executed
2. Your optimization will be based on **properties of the code that are true only in the common case you are considering**
  - ▶ If the code does not fit the common case, it shall stay as-is
  - ▶ Otherwise you **risk breaking program semantics!**

# Contents

Introduction

Compiler organization

Algorithm design

**Inside LLVM**

The LLVM-IR language

The Control Flow Graph

Conclusions

Bibliography

# LLVM-IR is like a box of chocolates

LLVM-IR comes in 3 different flavours:

**assembly** on-disk human-readable format  
(file extension: `.ll`)

**bitcode** on-disk machine-oriented binary format  
(file extension: `.bc`)

**in-memory** in-memory binary format  
(used during compilation process)

All formats have the same expressiveness!

# Using the driver to produce LLVM-IR

We can generate LLVM-IR assembly using the clang driver:

```
clang -emit-llvm -S -o out.ll in.c
```

If you want to generate bitcode instead:

```
clang -emit-llvm -o out.bc in.c
```

The compiler driver can also generate native code starting from  
LLVM-IR assembly

(Like compiling an assembly file with GCC)

# Playing with LLVM Passes

Run one or more passes on the LLVM-IR on-demand by using `opt`:

- ▶ Syntax is like `clang` (supports even `-O1`, `-O2`...)
- ▶ One command line argument per pass to run
- ▶ Order of execution is the same as the argument order
  - ▶ Different order, different results! (**phase/stage ordering**)

Some useful passes for debugging (they do not transform anything):

```
print CFG opt -view-cfg input.ll
```

```
print dominator tree opt -view-dom input.ll
```

```
print current IR opt -print-module input.ll
```

## Example

- ▶ Run *mem2reg*, then view the CFG:
  - ▶ `opt -mem2reg -view-cfg input.ll`

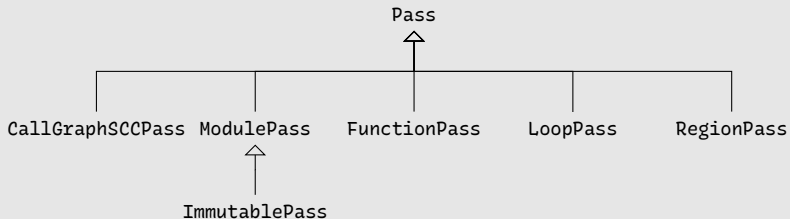
# Pass Hierarchy

LLVM provides a lot of passes...

- Try `opt -help!`

For performance reasons there are different kind of passes:

## LLVM Passes



# LLVM Passes

Each kind of pass visits particular elements of a module:

- ImmutablePass** compiler configuration – never run
- CallGraphSCCPass** post-order visit of CallGraph SCCs
- ModulePass** visit the whole module
- FunctionPass** visit functions
- LoopPass** post-order visit of loop nests
- RegionPass** visit a custom-defined region of code

Specializations come with restrictions:

- ▶ e.g. a **FunctionPass** cannot add or delete functions
- ▶ refer to “Writing a LLVM Pass” [4] for documentation on features and limitations of each kind of pass

# Recap

- ▶ The **user** invokes the **compiler** via the **driver**
- ▶ The **compiler** is made of three **stages**  
(front-end, middle-end, back-end)
- ▶ The **middle-end** is made of **passes**

If you want things done, you want to work on a **pass**.

- ▶ Edit an existing pass
- ▶ Create a new pass

To design a pass, you must follow the principle of **conservativeness**.

Next step: how to actually code a pass!



# Contents

Introduction

Compiler organization

Algorithm design

**Inside LLVM**

The LLVM-IR language

The Control Flow Graph

Conclusions

Bibliography

# How passes work

A **pass** is a **subroutine** that programmatically transforms a piece of code.

The code of the pass operates on the LLVM-IR using a set of **object-oriented APIs**.

Let's examine the LLVM-IR [5] more closely, first by looking at its **human-readable** form.

# LLVM-IR

```
define i32 @fact(i32 %n) {  
entry:  
    %retval = alloca i32, align 4  
    %n.addr = alloca i32, align 4  
    store i32 %n, i32* %n.addr, align 4  
    %0 = load i32, i32* %n.addr, align 4  
    %cmp = icmp eq i32 %0, 0  
    br i1 %cmp, label %if.then, label %if.end  
  
if.then:  
    store i32 1, i32* %retval, align 4  
    br label %return  
  
if.end:  
    %1 = load i32, i32* %n.addr, align 4  
    %2 = load i32, i32* %n.addr, align 4  
    %sub = sub nsw i32 %2, 1  
    %call = call i32 @fact(i32 %sub)  
    %mul = mul nsw i32 %1, %call  
    store i32 %mul, i32* %retval, align 4  
    br label %return  
  
return:  
    %3 = load i32, i32* %retval, align 4  
    ret i32 %3  
}
```

# LLVM-IR

LLVM-IR looks a lot like a RISC assembly language:

- ▶ Few instructions, all perfectly orthogonal
  - ▶ There are infinite registers
  - ▶ There are no special-purpose registers
  - ▶ No implicit flags register
- ▶ Basic block boundaries are denoted by **labels**
- ▶ Only **load** and **store** access memory

There are also a few CISC-like **high level instructions**:

- ▶ Reserve memory on the stack – **alloca**
- ▶ Function call – **call**
  - ▶ The calling convention is abstracted away
  - ▶ There is an implicit call stack
- ▶ Pointer arithmetics – **getelementptr**
- ▶ ...

# LLVM-IR

In reality LLVM-IR is much more high-level than assembly.

- ▶ The topmost object of a LLVM-IR program is the **Module**.
- ▶ **Modules** contain a list of **Globals**.
  - ▶ Globals can be either **Functions** or **Global Variables**.
  - ▶ A global can be a **Forward declaration**.
- ▶ **Functions** contain a list of **Basic Blocks** + **Arguments**.
- ▶ **Basic Blocks** are a list of **Instructions**.

## The in-memory representation

All these parts will correspond directly to **C++ objects**.

The abundance of lists guarantees low overhead and scalability to very large programs.

# LLVM-IR

LLVM-IR is **strongly typed**:

- ▶ e.g. you cannot assign a floating point value to an integer register without an explicit cast

**Almost everything is typed:**

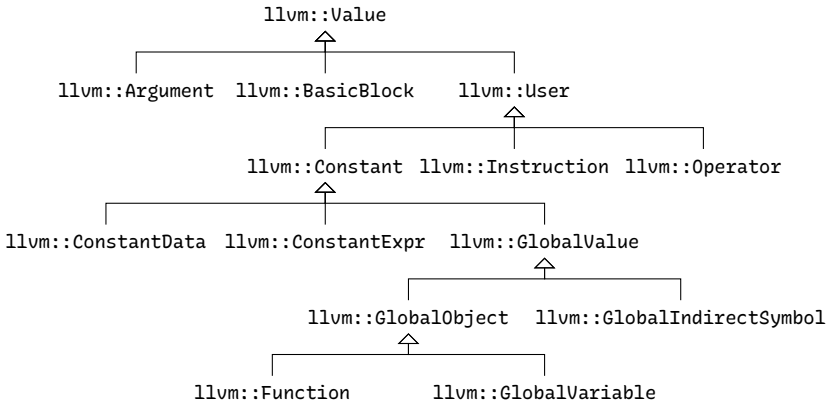
<b>functions</b>	@fact	→	i32 (i32)
<b>registers</b>	%3 = icmp eq i32 %2, 0	→	i1
<b>global vars.</b>	@var = common global i32 0	→	i32

These objects that have a type are called (somewhat confusingly) **LLVM Values**.

## The in-memory representation

`llvm::Value` is the **base class** of almost all interesting LLVM-IR objects!

# LLVM-IR



# LLVM-IR

LLVM-IR is SSA-based:

- ▶ every register is **statically assigned** exactly **once**

Statically means that:

- ▶ inside each function...
- ▶ ...for each variable `%foo`...
- ▶ ...there is **only one** statement in the form `%foo = ...`

**Static** (compile time)  $\neq$  **dynamic** (runtime)

- ▶ Single *Dynamic* Assignment:  
*in the execution trace* there is only one assignment to a variable `x`
- ▶ Single *Static* Assignment:  
*in the code listing* there is only one assignment to a variable `x`
  - ▶ Assignments **can** be performed multiple times (in a loop for example)



# Static Single Assignment

## Scalar SAXPY

```
float saxpy(float a, float x, float y) {  
    return a * x + y;  
}
```

## Scalar LLVM SAXPY

```
define float @saxpy(float %a, float %x, float %y) {  
    %1 = fmul float %a, %x  
    %2 = fadd float %1, %y  
    ret float %2  
}
```

Temporary %1 not reused! %2 is used for the second assignment!

# Static Single Assignment

## Array SAXPY

```
void saxpy(float a, float x[4], float y[4], float z[4]) {  
    for(unsigned i = 0; i < 4; ++i)  
        z[i] = a * x[i] + y[i];  
}
```

## Array LLVM SAXPY

```
for.cond:  
    %i.0 = phi i32 [ 0, %entry ], [ %inc, %for.inc ]  
    %cmp = icmp ult i32 %i.0, 4  
    br i1 %cmp, label %for.body, label %for.end  
[...]  
for.inc:  
    %inc = add i32 %i.0, 1  
    br label %for.cond
```

One assignment for loop counter %i.0

# Static Single Assignment

## Max

```
float max(float a, float b) {  
    return a > b ? a : b;  
}
```

## LLVM Max – WRONG

```
%1 = fcmp ogt float %a, %b  
br i1 %1, label %if.then, label %if.else  
if.then:  
    %2 = %a  
    br label %if.end  
if.else:  
    %2 = %b  
    br label %if.end  
if.end:  
    ret float %2
```

Why is it **wrong**?

# Static Single Assignment

The %2 variable must be statically assigned once!  
How do we handle conditional assignments then?

## LLVM Max

```
%1 = fcmp ogt float %a, %b
br i1 %1, label %if.then, label %if.end
if.then:
    br label %if.end
if.else:
    br label %if.end
if.end:
    %2 = phi float [ %a, %if.then ], [ %b, %if.else ]
    ret float %2
```

The **phi** instruction is a *conditional move*:

- ▶ it takes  $(\text{variable}_i, \text{label}_i)$  pairs
- ▶ if coming from predecessor identified by  $\text{label}_i$ , its value is  $\text{variable}_i$

# Static Single Assignment

Each SSA variable is assigned only once:

- ▶ variable **definition**

Each SSA variable can be referenced by multiple instructions:

- ▶ variable **uses**

Algorithms and technical language abuse of these terms!

*Let %foo be a variable. If the definition of %foo does not have side-effects nor uses, the aforementioned %foo variable can be erased from the CFG without altering program semantics.*

# SSA & LLVM-IR

## Important observation

SSA means that  
**there is always a 1:1 correspondence  
between a register and the instruction that assigns it.**

## Consequence

As a result, in LLVM-IR  
**registers are not separate objects  
but every LLVM Instruction  
is the output “register” of itself.**

# Static Single Assignment

Old compilers are not SSA-based:

- ▶ converting non-SSA input into SSA form is expensive
- ▶ cost must be amortized

New compilers are SSA-based:

- ▶ SSA easier to work with
- ▶ SSA-based analysis/optimizations are faster

# Contents

Introduction

Compiler organization

Algorithm design

**Inside LLVM**

The LLVM-IR language

The Control Flow Graph

Conclusions

Bibliography



# A step back...

Remember how we described the internal structure of an LLVM-IR module:

- ▶ `llvm::Module` is a list of `llvm::GlobalValues`.
  - ▶ `llvm::Function` is a kind of `llvm::GlobalValue`.
- ▶ `llvm::Function` is a list of `llvm::BasicBlocks`.
- ▶ `llvm::BasicBlock` is a list of `llvm::Instructions`.

Functions and basic blocks act like containers:

- ▶ STL-like accessors: `front()`, `back()`, `size()`, ...
- ▶ STL-like iterators: `begin()`, `end()`

Each contained element is aware of its container:

- ▶ `getParent()`

Warning for BBs: order of iteration  $\neq$  order of execution!

# A step back...

In a `llvm::BasicBlock`, the `llvm::Instructions` execute in the order specified by the list.

- In which order do the `llvm::BasicBlocks` execute?

The way the basic blocks are executed is implicitly described by the **branches** in each block.

- These branches describe the **Control Flow Graph** of the function.

# Control Flow Graph

LLVM automatically maintains a simple API for operating on the CFG:

- ▶ no need to run passes
- ▶ no need to search the branch instructions in each basic block

Every CFG has an **entry** basic block:

- ▶ the **first** executed basic block
- ▶ it is the **root/source** of the graph
- ▶ get it with `llvm::Function::getEntryBlock()`

# Control Flow Graph

At the end of a basic blocks there's always a **terminator** instruction:

- ▶ **ret, br, switch, unreachable, ...**

More than one **exit** block can be present in a function:

- ▶ they are the **leaves/sinks** of the graph
- ▶ their terminator instructions are always **rets**
  1. `llvm::BasicBlock::getTerminator()`
  2. check the opcode of the terminator

# Side Note

For performance reasons, a custom casting framework is used:

- ▶ you cannot use **static\_cast** and **dynamic\_cast** with types/classes provided by LLVM

## LLVM Casting Functions

Static cast of Y* to X	<code>X *llvm::cast&lt;X&gt;(Y *)</code>
Dynamic cast of Y* to X	<code>X *llvm::dyn_cast&lt;X&gt;(Y *)</code>
Is Y* an instance of X?	<code>bool llvm::isa&lt;X&gt;(Y *)</code>

Example:

- ▶ is BB a sink?  
`llvm::isa<llvm::ReturnInst>(BB.getTerminator())`

# Control Flow Graph

Every basic block BB has one or more\*:

**predecessors** from `pred_begin(BB)` to `pred_end(BB)`

**successors** from `succ_begin(BB)` to `succ_end(BB)`

Other convenience methods available in `llvm::BasicBlock`:

- ▶ useful getters
  - ▶ `BasicBlock *getUniquePredecessor()`
  - ▶ ...
- ▶ moving a basic block
  - ▶ `moveBefore(llvm::BasicBlock *)`
  - ▶ `moveAfter(llvm::BasicBlock *)`
- ▶ split a basic block:
  - ▶ `splitBasicBlock(llvm::BasicBlock::iterator)`
- ▶ ...

---

\*see `include/llvm/IR/CFG.h`

# Control Flow Graph

The `llvm::Instruction` class defines common operations:

- ▶ getting an operand
  - ▶ `getOperand(unsigned)`

Subclasses provide specialized accessors:

- ▶ the **load** instruction takes as operand the pointer to the memory to be loaded:
  - ▶ `llvm::LoadInst::getPointerOperand()`

# Instructions

Instructions are created using:

- ▶ constructors
  - ▶ `llvm::LoadInst::LoadInst(...)`
- ▶ factory methods
  - ▶ `llvm::GetElementPtrInst::Create(...)`
- ▶ the helper class `llvm::IRBuilder`
  - ▶ `llvm::IRBuilder<> builder(insPoint);`  
`builder.CreateAdd(...);`

**Interface is not homogeneous!**

Some instructions support all methods, others support only one.



# Instructions

Instructions can be inserted:

- ▶ automatically by `IRBuilder`
  - ▶ insertion point is given at `IRBuilder` instantiation
- ▶ manually by appending to a basic block
- ▶ manually by inserting after/before another instruction

# From Control Flow to Data Flow

In LLVM, the data flow generated by the various instructions is represented by a simple hierarchy:

**value** something that can be used: `llvm::Value`

**user** something that can use: `llvm::User`

**use** the link between the **value** and the **user**: `llvm::Use`

A value is a **definition**:

- ▶ Visiting where a definition is used:
  - ▶ `llvm::Value::use_begin()`, `llvm::Value::use_end()`

An user accesses **definitions**:

- ▶ Visiting the definitions that are used:
  - ▶ `llvm::User::op_begin()`, `llvm::User::op_end()`

# From Control Flow to Data Flow

- ▶ `llvm::Value` inherits from `llvm::User`
- ▶ `llvm::Instruction` inherits from `llvm::Value`
  - ⇒ The value produced by the instruction is the **instruction itself**!

## Example

```
%6 = load i32, i32* %1, align 4
```

The **load** is described by an instance of `llvm::Instruction`.  
That instance also represents the `%6` variable.

Not all instances of `llvm::Value` are also `llvm::Instructions`!  
i.e. function arguments

# From Control Flow to Data Flow

Every `llvm::Value` is typed:

- ▶ use `llvm::Value::getType()` to get the type

Since every instruction is a value:

- ▶ instructions are typed

## Example

```
%6 = load i32, i32* %1, align 4
```

The type of the `%6` variable is the type of the return value of the **load** instruction, `i32`

# Contents

Introduction

Compiler organization

Algorithm design

Inside LLVM

**Conclusions**

Bibliography

# Conclusions

LLVM is a **production-quality** compiler framework:

⇒ impossible knowing all details

But:

- ▶ it is well organized
- ▶ if you know compiler theory, it is relatively easy to find what you need inside the source code

Remember it's written in C++!

- ▶ To hack around LLVM you need at least basic C++ skills
- ▶ C++  $\neq$  C

# Thank You!

Questions?

# Contents

Introduction

Compiler organization

Algorithm design

Inside LLVM

Conclusions

**Bibliography**



# Bibliography I



University of Illinois at Urbana-Champaign.  
Low Level Virtual Machine.  
<http://www.llvm.org>.



GNU.  
GNU Compiler Collection.  
<http://gcc.gnu.org>.



University of Illinois at Urbana-Champaign.  
Clang: a C language family frontend for LLVM.  
<http://clang.llvm.org>.



Chris Lattner and Jim Laskey.  
Writing an LLVM Pass.  
<http://llvm.org/docs/WritingAnLLVMPass.html>.

# Bibliography II



Chris Lattner and Vikram Adve.  
LLVM Language Reference Manual.  
<http://llvm.org/docs/LangRef.html>.



Linus Torvalds.  
Re: SCO: "thread creation is about a thousand times faster than  
onnative.  
<https://lkml.org/lkml/2000/8/25/132>.



Bruce Eckel.  
Thinking in C++ – Volume One: Introduction to Standard C++.  
<http://mindview.net/Books/TICPP/ThinkingInCPP2e.html>.



Bruce Eckel and Chuck Allison.  
Thinking in C++ – Volume Two: Practical Programming.  
<http://mindview.net/Books/TICPP/ThinkingInCPP2e.html>.

# Bibliography III



AMD.

Open64.

<http://developer.amd.com/tools-and-sdks/cpu-development/x86-open64-compiler-suite>.



John T. Criswell, Daniel Dunbar, Reid Spencer, and Tanya Lattner.

LLVM Testing Infrastructure Guide.

<http://llvm.org/docs/TestingGuide.html>.



LLVM Community.

LLVM Coding Standards.

<http://llvm.org/docs/CodingStandards.html>.



LLVM Community.

LLVM Passes.

<http://llvm.org/docs/Passes.html>.

# Bibliography IV



LLVM Community.  
Autovectorization in LLVM.  
<http://llvm.org/docs/Vectorizers.html>.



LLVM Community.  
LLVM Programmer's Manual.  
<http://llvm.org/docs/ProgrammersManual.html>.



Ettore Speziale.  
Compiler Optimization and Transformation Passes.  
<https://github.com/speziale-ettore/COTPasses>.



Scott Chacon.  
Pro Git.  
<http://git-scm.com/book>.