

# Introduction to LLVM compiler framework

Stefano Cherubin

Politecnico di Milano

03-05-2019

*This material is strongly based on material produced by Michele Scandale and Ettore Speziale for the course 'Code Optimizations and Transformations'.*

# Contents

**Introduction**

Compiler organization

Algorithm design

Inside LLVM

LLVM-IR language

Conclusions

Bibliography

# Compilers and compilers

Approaching to compilers, we need to understand the difference between a *toy-compiler* and *production-quality compiler*.

## Toy Compiler

- ▶ small code-base
- ▶ easy doing tiny edits
- ▶ impossible doing normal/big edits

## Production-Quality Compiler

- ▶ huge code-base
- ▶ difficult performing any kind of edits
- ▶ compiler-code extremely optimized

Key concepts:

- ▶ working with a production-quality compiler is *initially hard*, but . . .
- ▶ . . . an huge set of tools for analyzing/transforming/testing code is provided – toy compilers **miss these things!**

# LLVM: Low Level Virtual Machine

Initially started as a research project at Urbana-Champaign:

- ▶ now intensively used for **researches** involving compilers
- ▶ key technology for **leading industries** – AMD, Apple, Intel, NVIDIA

If you are there, then it is **your key-technology**:

- ▶ open-source compilers: GCC [?], LLVM [?]
- ▶ LLVM is relatively **young** – GCC performances may be better –  
...
- ▶ ... LLVM is more modular, well written, kept *clean* by developers.

# Contents

Introduction

**Compiler organization**

Algorithm design

Inside LLVM


LLVM-IR language

Conclusions

Bibliography

# Compiler pipeline

Typically a compiler is a **pipeline**:



There are three main components:

**Front-end** **translate** a source file into the intermediate representation

**Middle-end** **analyze** intermediate representation, **optimize** it

**Back-end** **generate** target machine assembly from the intermediate representation

# Compiler pipeline

Each component is composed internally by pipelines:

- ▶ simple model – read something, produce something
- ▶ specify only how to transform input data into output data

Complexity lies on **chaining** together stages.

# Compiler pipeline

We will consider only the *middle-end*: same concepts are valid also for {front,back}-end.

Technical terms:

**Pass** a pipeline stage

**IR** (a.k.a. Intermediate Representation) is the language used in the middle-end.

The **pass manager** manages a set of passes:

- build the compilation pipeline: **schedule** passes together according to **dependencies**.

Dependencies are **hints** used by the pass manager in order to schedule passes.



# First insights

A compiler is **complex**:

- ▶ passes are the **elementary unit of work**
- ▶ pass manager must be **advisee** about pass chaining
- ▶ pipeline shapes are **not fixed** – it can change from one compiler execution to another <sup>1</sup>

Moreover, compilers must be **conservative**:

- ▶ apply a transformation only if program **semantic is preserved**

Compiler algorithms are designed differently w.r.t. standard algorithms!

---

<sup>1</sup>e.g. optimized/not optimized builds, compiler options, ...

# Contents

Introduction

Compiler organization

**Algorithm design**

Inside LLVM

LLVM-IR language

Conclusions

Bibliography

# Classical Algorithm Design

Dealing with algorithm design, a good approach is the following:

1. study the problem
2. make some example
3. identify the **common case**
4. derive the algorithm for the common case
5. add handling for **corner cases**
6. improve performing **optimizing the common case**

Weakness of the approach:

- ▶ **corner cases** – a *correct* algorithm **must** consider *all the corner cases*!

# Compiler Algorithm Design

Corner cases are difficult to handle:

- ▶ compiler algorithms must be **proved** to preserve program semantic
- ▶ having a common methodology helps on that

Compiler algorithms are built combining three kind of **passes**:

- ▶ analysis
- ▶ optimization
- ▶ (normalization)

# Compiler Algorithm Design

Corner cases are difficult to handle:

- ▶ compiler algorithms must be **proved** to preserve program semantic
- ▶ having a common methodology helps on that

Compiler algorithms are built combining three kind of **passes**:

- ▶ analysis
- ▶ optimization
- ▶ (normalization)

We now consider a simple example: *loop hoisting*.

# Loop Hoisting

It is a transformation that:

- ▶ looks for statements (inside a loop) not depending on the loop state
- ▶ move them outside the loop body

## Loop Hoisting – Before

```
do {  
    a += i;  
    b = c;  
    i++;  
} while (i < k);
```

## Loop Hoisting – After

```
b = c;  
do {  
    a += i;  
    i++;  
} while (i < k);
```

# Loop Hoisting

## Transformation

The transformation is trivial:

- ▶ move “good” statement outside of the loop

This is the **optimization pass**. It needs to know:

- ▶ which pieces of code are loops
- ▶ which statements are “good” statements

They are **analysis**, which have to be implemented by other passes:

- ▶ detecting loops in the program
- ▶ detecting loop-independent statements

When registering loop hoisting, also declare needed analysis:

- ▶ pipeline automatically built: **analysis** → **optimization**

# Loop Hoisting

The **proof** is trivial:

- ▶ transformation is correct if analysis are correct, but . . .
- ▶ . . . usually analysis are built starting from other analysis already implemented inside the compiler

You have to prove that combining all analysis information gives you a correct view of the code:

- ▶ analysis information cannot induce optimization passes applying a transformation not preserving program semantic



# Loop Hoisting

We have spoken about loops, but which kind of loop?

- ▶ do-while loops?
- ▶ while loop?
- ▶ for loops?

We have seen loop hoisting on:

- ▶ do-while loops

What about other kinds of loops?

- ▶ they must be normalized – i.e. transformed to do-while loops

**Normalization passes** do that:

- ▶ before running loop hoisting, you must tell to the pass manager that loop normalization must be run before

This allows to recognize more loops, thus potentially **improving optimization impact!**

# Compiler Algorithm Design

You have to:

1. analyze the problem
2. make some examples
3. detect the common case
4. declare the **input format**
5. declare **analysis** you need
6. design an **optimization** pass
7. proof its **correctness**
8. improve algorithm performance by acting on common case – the only considered up to now. Please notice that corner cases are not considered – just do not try to optimize the corner cases
9. improve the effectiveness of the algorithm by adding **normalization passes**

# Contents

Introduction

Compiler organization

Algorithm design

**Inside LLVM**

LLVM-IR language

Conclusions

Bibliography

# Terminology

LLVM IR comes with 3 different flavours:

**assembly** human-readable format

**bitcode** binary on-disk machine-oriented format

**in-memory** binary in-memory format, used during compilation process

All formats have the same expressiveness!

File extensions:

**.ll** for assembly files

**.bc** for bitcode files

# Tools

Writing LLVM assembly by hand is unfeasible:

- ▶ different front-ends available for LLVM
- ▶ use Clang [?] for the C family

The clang driver is compatible with GCC:

- ▶  $\approx$  same command line options

To generate LLVM IR:

**assembly** `clang -emit-llvm -S -o out.ll in.c`

**bitcode** `clang -emit-llvm -o out.bc in.c`

It can also generate native code starting from LLVM assembly or LLVM bitcode – like compiling an assembly file with GCC

# Tools

LLVM IR can be manipulated using `opt`:

- ▶ read an input file
- ▶ run specified LLVM passes on it
- ▶ respecting user-provided order

Useful passes:

- ▶ print CFG with `opt -view-cfg input.ll`
- ▶ print dominator tree with `opt -view-dom input.ll`
- ▶ ...

Pass chaining:

- ▶ run *mem2reg*, then view the CFG with  
`opt -mem2reg -view-cfg input.ll`
- ▶ potentially different results using different option order  
**(phase/stage ordering)**


# Pass Hierarchy

LLVM provides a lot of passes:

► `try opt -help`

For performance reasons there are different kind of passes:

## LLVM Passes



# LLVM Passes

Each pass kind visits particular elements of a module:

- ImmutablePass** compiler configuration – never run
- CallGraphSCCPass** post-order visit of CallGraph SCCs
- ModulePass** visit the whole module
- FunctionPass** visit functions
- LoopPass** post-order visit of loop nests
- RegionPass** visit a custom-defined region of code

Specializations comes with restrictions:

- ▶ e.g. a **FunctionPass** cannot add or delete functions
- ▶ refer to “Writing a LLVM Pass” [?] for accurate description of features and limitations of each kind of pass



# What is Available Inside LLVM?

LLVM provides passes performing basic transformations:

- ▶ variables promotion
- ▶ loops canonicalization
- ▶ ...

They can be used to **normalize/canonicalize** the input

- ▶ transform into a form analyzable for further passes
- ▶ it is essential because keeps passes implementation manageable

# Contents

Introduction

Compiler organization

Algorithm design

Inside LLVM

**LLVM-IR language**

Conclusions

Bibliography

# LLVM IR

LLVM IR [?] language is RISC-based:

- ▶ instructions operates on **variables** <sup>2</sup>
- ▶ only load and store access memory
- ▶ `alloca` used to reserve memory on function stacks

There are also few **high level instructions**:

- ▶ function call – `call`
- ▶ pointer arithmetics – `getelementptr`
- ▶ ...

---

<sup>2</sup>Virtual registers

# LLVM IR

LLVM IR is **strongly typed**:

- ▶ e.g. you cannot assign a floating point value to an integer variable without an explicit cast

**Almost everything is typed** – e.g.:

**functions** @fact – i32 (i32)

**statements** %3 = icmp eq i32 %2, 0 – i1

A variable can be:

**global** @var = common global i32 0, align 4

**function parameter** define i32 @fact(i32 %n)

**local** %2 = load i32, i32\* %1, align 4

Local variables are defined by statements

# LLVM IR

```
define i32 @fact(i32 %n) {  
entry:  
    %retval = alloca i32, align 4  
    %n.addr = alloca i32, align 4  
    store i32 %n, i32* %n.addr, align 4  
    %0 = load i32, i32* %n.addr, align 4  
    %cmp = icmp eq i32 %0, 0  
    br i1 %cmp, label %if.then, label %if.end  
  
if.then:  
    store i32 1, i32* %retval, align 4  
    br label %return  
  
if.end:  
    %1 = load i32, i32* %n.addr, align 4  
    %2 = load i32, i32* %n.addr, align 4  
    %sub = sub nsw i32 %2, 1  
    %1 = mul i32 %1, %sub  
    store i32 %1, i32* %retval, align 4  
    br label %if.end  
}
```

# LLVM IR Language

LLVM IR is SSA-based:

- ▶ every variable is **statically assigned** exactly **once**

Statically means that:

- ▶ inside each function
- ▶ for each variable `%foo`
- ▶ there is only one statement in the form `%foo = ...`

Static is different from dynamic:

- ▶ a static assignment can be executed more than once

# Static Single Assignment

## Scalar SAXPY

```
float saxpy(float a, float x, float y) {  
    return a * x + y;  
}
```

## Scalar LLVM SAXPY

```
define float @saxpy(float %a, float %x, float %y) {  
    %1 = fmul float %a, %x  
    %2 = fadd float %1, %y  
    ret float %2  
}
```

Temporary %1 not reused! %2 is used for the second assignment!

# Static Single Assignment

## Array SAXPY

```
void saxpy(float a, float x[4], float y[4], float z[4])  
    for(unsigned i = 0; i < 4; ++i)  
        z[i] = a * x[i] + y[i];  
}
```

## Array LLVM SAXPY

```
for.cond:  
    %i.0 = phi i32 [ 0, %entry ], [ %inc, %for.inc ]  
    %cmp = icmp ult i32 %i.0, 4  
    br i1 %cmp, label %for.body, label %for.end  
  
...  
  
for.inc:
```



# Static Single Assignment

## Max

```
float max(float a, float b) {  
    return a > b ? a : b;  
}
```

## LLVM Max – Bad

```
%1 = fcmp ogt float %a, %b  
br i1 %1, label %if.then, label %if.else  
if.then:  
    %2 = %a  
    br label %if.end  
if.else:  
    %2 = %b  
    br label %if.end  
if.end:
```

# Static Single Assignment

The %2 variable must be statically set once

## LLVM Max

```
%1 = fcmp ogt float %a, %b
br i1 %1, label %if.then, label %if.end
if.then:
    br label %if.end
if.else:
    br label %if.end
if.end:
    %2 = phi float [ %a, %if.then ], [ %b, %if.else ]
    ret float %2
```

The phi instruction is a *conditional move*:

- ▶ it takes  $(\text{variable}_i, \text{label}_i)$  pairs
- ▶ if coming from predecessor identified by  $\text{label}_i$ , its value is  $\text{variable}_i$

# Static Single Assignment

Each SSA variable is set only once:

- ▶ variable **definition**

Each SSA variable can be used by multiple instructions:

- ▶ variable **uses**

Algorithms and technical language abuse of these terms:

*Let  $\%foo$  be a variable. If  $\%foo$  definition has not side-effects, and no uses, dead-code elimination can be efficiently performed by erasing  $\%foo$  definition from the CFG.*

# Static Single Assignment

Old compilers are not SSA-based:

- ▶ putting input into SSA-form is expensive
- ▶ cost must be amortized

New compilers are SSA-based:

- ▶ SSA easier to work with
- ▶ SSA-based analysis/optimizations faster

# Contents

Introduction

Compiler organization

Algorithm design

Inside LLVM

LLVM-IR language

**Conclusions**

Bibliography

# Conclusions

LLVM is a **production-quality** compiler framework:

⇒ impossible knowing all details

But:

- ▶ it is well organized
- ▶ given you known compilers theory, it is relatively easy to find what you need inside its sources

Please take into account C++:

- ▶ basic skills required

# Contents

Introduction

Compiler organization

Algorithm design

Inside LLVM

LLVM-IR language

Conclusions

**Bibliography**

# **Bibliography I**