

# Exploring LLVM

Stefano Cherubin

Politecnico di Milano

08-05-2019

*This material is strongly based on material produced by Michele Scandale and Ettore Speziale for the course 'Code Optimizations and Transformations'.*

# Contents

# LLVM official documentation

[llvm.org/docs](http://llvm.org/docs)

# A lot of documentation...

[llvm.org/docs](http://llvm.org/docs) mentions:

- 5 references about *Design & Overview*
- 23 references about *User Guides*
- 15 references about *Programming Documentation*
- 42 references about *Subsystem Documentation*
- 9 references about *Development Process Documentation*
- 5 Mailing Lists
- 5 IRC bots

Most of the above references are OUT-OF-DATE.

You probably need documentation about the documentation itself.

# Essential documentation

[Intro to LLVM](#) [?] gives a quick and clear introduction to the compiler infrastructure. It is mostly up-to-date.<sup>1</sup>

[Writing an LLVM pass](#) [?] explains step by step how to implement a Pass for those who never did anything like that. We will see this tutorial later in the course.

[Doxygen](#) [?] *The best code documentation is the code itself.* Sometimes the generated doxygen documentation is enough. It also contains links to the web version of the source code. It is updated to the latest development branch. Please refer to github branches for the documentation about the stable versions.

[llvm-dev](#) Mailing List. Last resource: ask other developers.  
Warning: 24/7 many people are posting in this ML.

---

<sup>1</sup>at the time I am writing

# Contents

# Canonicalize Pass Input

We will see the following passes:

Pass	Switch
Variable promotion	mem2reg
Loop simplify	loop-simplify
Loop-closed SSA	lcssa
Induction variable simplification	indvars

They are **normalization** passes:

- put data into a canonical form

# Variable Promotion

One of the most difficult things in compiler is:

- considering memory accesses

## Plain SAXPY

```
define float @saxpy(float %a, float %x, float %y) {  
entry:  
    %a.addr = alloca float, align 4  
    %x.addr = alloca float, align 4  
    %y.addr = alloca float, align 4  
    store float %a, float* %a.addr, align 4  
    store float %x, float* %x.addr, align 4  
    store float %y, float* %y.addr, align 4  
    %0 = load float, float* %a.addr, align 4  
    %1 = load float, float* %x.addr, align 4  
    %mul = fmul float %0, %1  
    %2 = load float, float* %y.addr, align 4  
    %add = fadd float %mul, %2  
    ret float %add  
}
```



# Variable Promotion

## Simplifying Representation

In the SAXPY kernel some **alloca** are generated:

- represent **local variables**<sup>2</sup>

They are generated due to compiler **conservative** approach:

- maybe some instruction can take the addresses of such variables, hence a memory location is needed

Complex representations makes hard performing further actions:

- suppose you want to compute  $a * x + y$  using only one instruction<sup>3</sup>
- hard to detect due to **load** and **store**

---

<sup>2</sup>Arguments are considered local variables

<sup>3</sup>e.g. FMA4

# Variable Promotion

## Using Memory Only When Necessary

To limit the number of instruction accessing memory:

- we need to eliminate **load** and **store**
- achieved by **promoting** variables from memory to registers

Inside LLVM SSA-based representation:

**memory** Stack allocations – e.g. `%1 = alloca float, align 4`

**register** SSA variables – e.g. `%a`

The **mem2reg** pass focus on:

- eliminating **alloca** with only **load** and **store** uses

Also available as utility:

- `llvm::PromoteMemToReg`<sup>4</sup>

---

<sup>4</sup>see `lib/Transforms/Utils/PromoteMemoryToRegister.cpp`

# Variable Promotion

Example on simplified code

## Starting Point

```
%1 = alloca float
%2 = alloca float
%3 = alloca float
store %a, %1
store %x, %2
store %y, %3
%4 = load %1
%5 = load %2
%6 = fmul %4, %5
%7 = load %3
%8 = fadd %6, %7
ret %8
```

Copy propagation performed  
transparently by the compiler

## Promoting `alloca`

```
%1 = %a
%2 = %x
%3 = %y
%4 = %1
%5 = %2
%6 = fmul %4, %5
%7 = %3
%8 = fadd %6, %7
ret %8
```

## After Copy-propagation

```
%1 = fmul %a, %x
%2 = fadd %1, %y
ret %2
```


# Loops

Different kind of loops:


**do-while** Loops



**while** Loops



**Irreducible** Loops



In LLVM the focus is on one kind of loop:

- natural loops

# Natural Loops

A natural loop:

- has only one entry node – *header*
- there is a back edge that enter the loop header

Under this definition:

- the irreducible loop is not a natural loop
- since LLVM consider only natural loops, the irreducible loop **is not recognized** as a loop

# Loop Terminology

Loops defined starting from back-edges:


**back-edge** edge entering loop header: (3, 1)

**header** loop entry node: 1

**body** nodes that can reach  
back-edge source node (3)  
without passing from  
back-edge target node (1)  
plus back-edge target node:  
{1, 2, 3}

**exiting** nodes with a successor outside the loop: {1, 3}

**exit** nodes with a predecessor inside the loop: {4, 5}



# Loop Simplify

Natural loops finding is the base pass **identify** loops, but:

- some features are not analysis/optimization friendly


The **loop-simplify** pass normalize natural loops:

**pre-header** the **only predecessor** of **header** node

**latch** the **starting node** of the **only back-edge**

**exit-block** ensures **exits** **dominated** by loop **header**


## Pre-header Insertion




# Loop Simplify

## Example

### Latch Insertion



### Exit-block Insertion



- pre-header always executed before entering the loop
- latch always executed before starting a new iteration
- exit-blocks always executed after exiting the loop



# Loop-closed SSA

Loop representation can be further normalized:

- **loop-simplify** normalize the **shape** of the loop
- nothing is said about loop definitions

Keeping SSA form is expensive with loops:

- **lcssa** insert **phi** instruction at loop boundaries for variables **defined inside** the loop body and **used outside**
- this guarantees isolation between optimization performed inside and outside the loop
- faster keeping IR into SSA form – propagation of code changes outside the loop blocked by **phi** instructions

# Loop-closed SSA

## Example

### Linear Search

```
unsigned search(float *x, unsigned n, float y) {  
    unsigned i, j = 0;  
    for(i = 0; i != n; ++i)  
        if(x[i] == y)  
            j = i;  
    return j;  
}
```

The example is trivial:

- think about having large loop bodies
- transformation becomes useful

# Loop-closed SSA

## Example

### Before LCSSA

```
for.cond:  
  %i.0 = phi i32 [ 0, %entry ], [ %inc, %for.inc ]  
  %j.0 = phi i32 [ 0, %entry ], [ %j.1, %for.inc ]  
  %cmp = icmp ne i32 %i.0, %n  
  br i1 %cmp, label %for.body, label %for.end  
  
  ...  
  
if.end:  
  %j.1 = phi i32 [ %i.0, %if.then ], [ %j.0, %for.body ]  
  br label %for.inc  
  
for.inc:  
  %inc = add i32 %i.0, 1  
  br label %for.cond  
  
for.end:  
  ret i32 %j.0
```

# Loop-closed SSA

## Example

### After LCSSA

for.cond:

**%i.0** = **phi** i32 [ 0, %entry ], [ %inc, %for.inc ]

**%j.0** = **phi** i32 [ 0, %entry ], [ %j.1, %for.inc ]

**%cmp** = **icmp ne** i32 %i.0, %n

**br** i1 %cmp, label %for.body, label %for.end

...

if.end:

**%j.1** = **phi** i32 [ %i.0, %if.then ], [ %j.0, %for.body ]

**br** label %for.inc

for.inc:

**%inc** = **add** i32 %i.0, 1

**br** label %for.cond

for.end:

**%j.0.lcssa** = **phi** i32 [ %j.0, %for.cond ]

**ret** i32 %j.0.lcssa

# Induction Variables

Some loop variables are *special*:

- e.g. counters

Generalization lead to **induction variables**:

- `foo` is a loop induction variable if its successive values form an arithmetic progression:

$$\text{foo} = \text{bar} * \text{baz} + \text{biz}$$

where `bar`, `biz` are loop-invariant<sup>5</sup>, and `baz` is an induction variable

- `foo` is a **canonical** induction variable if it is always incremented by a constant amount:

$$\text{foo} = \text{foo} + \text{biz}$$

where `biz` is loop-invariant

---

<sup>5</sup>Constants inside the loop

# Induction Variable Simplification

Canonical induction variables are used to **drive** loop execution:

- given a loop, the **indvars** pass tries to find its canonical induction variable

With respect to theory, LLVM canonical induction variable is:

- initialized to 0
- incremented by 1 at each loop iteration

# Normalization

## Wrap-up

Normalization passes running order:

- 1 **mem2reg**: limit use of memory, increasing the effectiveness of subsequent passes
- 2 **loop-simplify**: canonicalize loop shape, lower burden of writing passes
- 3 **lcssa**: keep effects of subsequent loop optimizations local, limiting overhead of maintaining SSA form
- 4 **indvars**: normalize induction variables, highlighting the canonical induction variable

Other normalization passes available:

- try running **opt -help**

# Contents



# Checking Input Properties

Analysis basically allows to:

- **derive** information and properties of the input
- **verify** properties of input

Keeping analysis information is expensive:

- tuned algorithms updates analysis information when an optimization invalidates them
- incrementally updating analysis is cheaper than recomputing them

Many LLVM analysis supports incremental updates:

- this is an **optimization**
- focus on **information** provided by analysis

# Useful Analysis

We will see the following passes:

## Analysis

Pass	Switch	Transitive
Control flow graph	<code>none</code>	No
Dominator tree	<code>domtree</code>	No
Post-dominator tree	<code>postdomtree</code>	No
Loop information	<code>loops</code>	Yes
Scalar evolution	<code>scalar-evolution</code>	Yes
Alias analysis	<code>special</code>	Yes
Memory dependence	<code>memdep</code>	Yes

# Require Analysis

Ask the pass manager to schedule a specific pass before running the current one.

Requiring analysis by transitivity:

**yes** `llvm::AnalysisUsage::addRequiredTransitive<T>()`

**no** `llvm::AnalysisUsage::addRequired<T>()`

In cases where **analyses chain**, the `addRequiredTransitive` method should be used instead of the `addRequired` method.

This informs the `PassManager` that the transitively required pass should be alive as long as the requiring pass is.

# Control Flow Graph

The Control Flow Graph is implicitly maintained by LLVM:

- no specific pass to build it

Recap:

- CFG for a function is a set of basic blocks
- a basic block is a set of instructions

Functions and basic blocks acts like containers:

- STL-like accessors: `front()`, `back()`, `size()`, ...
- STL-like iterators: `begin()`, `end()`

Each contained element is aware of its container:

- `getParent()`

# Control Flow Graph

## Walking

Every CFG has an entry basic block:

- the **first** executed basic block
- it is the **root/source** of the graph
- get it with `llvm::Function::getEntryBlock()`

More than one exit blocks can be generated:

- their terminator instructions are **rets**
- they are the **leaves/sinks** of the graph
- USE `llvm::BasicBlock::getTerminator()` to get the terminator ...
- ... then check its real class

# Side Note

## Casting Framework

For performance reasons, a custom casting framework is used:

- you cannot use `static_cast` and `dynamic_cast` with types/classes provided by LLVM

### LLVM Casting Functions

Meaning	Function
Static cast of <code>Y *</code> to <code>X *</code>	<code>X * llvm::cast&lt;X&gt;(Y *)</code>
Dynamic cast of <code>Y *</code> to <code>X *</code>	<code>X * llvm::dyn_cast&lt;X&gt;(Y *)</code>
Is <code>Y</code> an <code>X</code> ?	<code>bool llvm::isa&lt;X&gt;(Y *)</code>

Example:

- is `BB` a sink?

```
llvm::isa<llvm::ReturnInst>(BB.getTerminator())
```

# Control Flow Graph

## Basic Blocks

Every basic block `BB` has one or more:

`predecessors` from `pred_begin(BB)` to `pred_end(BB)` <sup>6</sup>

`successors` from `succ_begin(BB)` to `succ_end(BB)`

Convenience accessors directly available in `llvm::BasicBlock`:

- e.g. `llvm::BasicBlock::getUniquePredecessor()`

Other convenience member functions:

- moving a basic block: `llvm::BasicBlock::moveBefore(llvm::BasicBlock *)` OR `llvm::BasicBlock::moveAfter(llvm::BasicBlock *)`
- split a basic block:  
`llvm::BasicBlock::splitBasicBlock(llvm::BasicBlock::iterator)`
- ...

---

<sup>6</sup>see `include/llvm/IR/CFG.h`

# Control Flow Graph

## Instructions

The `llvm::Instruction` class define common operations:

- e.g. getting an operand: `llvm::Instruction::getOperand(unsigned)`

Subclasses provide specialized accessors:

- e.g the `load` instruction takes an operand that is a pointer:

`llvm::LoadInst::getPointerOperand()`



# Control Flow Graph

## Instructions

The `llvm::Instruction` class define common operations:

- e.g. getting an operand: `llvm::Instruction::getOperand(unsigned)`

Subclasses provide specialized accessors:

- e.g the `load` instruction takes an operand that is a pointer:

`llvm::LoadInst::getPointerOperand()`

The value produced by the instruction is the **instruction itself**:

### Example

Consider:

```
%6 = load i32, i32* %1, align 4
```

the `load` is described by an instance of `llvm::LoadInst`. That instance also models the `%6` variable

# Instructions

## Creating New Instructions

Instructions built using:

- constructors – e.g. `llvm::LoadInst::LoadInst(...)`
- factory methods – e.g. `llvm::GetElementPtrInst::Create(...)`

Interface is not homogeneous:

- some instructions support both methods
- others support only one

At build-time, instructions can be:

- appended to a basic block
- inserted after/before a given instruction

Insertion point usually specified as builder last argument

# Side Note

## Definitions and Uses

LLVM class hierarchy is built around two simple concepts:

**value** something that can be used: `llvm::Value`

**user** something that can use: `llvm::User`

A value is a **definition**:

- `llvm::Value::use_begin()`, `llvm::Value::use_end()` to visit uses <sup>7</sup>

An user access **definitions**:

- `llvm::User::op_begin()`, `llvm::User::op_end()` to visit used values <sup>8</sup>

Functions:

- used by call sites
- uses formal parameters

Instructions:

- define an SSA value
- uses operands

---

<sup>7</sup>`llvm::Instruction` derives from `llvm::Value`

<sup>8</sup>`llvm::Value` derives from `llvm::User`

# Side Note

## Value Typing

Every `llvm::Value` is typed:

- USE `llvm::Value::getType()` to get the type

Since every instructions is/define a value:

- instructions are typed

### Example

Consider:

```
%6 = load i32, i32* %1, align 4
```

- The `%6` variable actually is the instruction itself
- Its type is the type of `load` return value, `i32`

# Dominance Trees

Dominance trees answer to control-related queries:

- is this basic block executed before that?
- `llvm::DominatorTree`
- is this basic block executed after that?
- `llvm::PostDominatorTree`

The two trees interface is similar:

- `bool dominates(X *, X *)`
- `bool properlyDominates(X *, X *)`

Where `x` is an `llvm::BasicBlock` OR an `llvm::Instruction`

by using `opt`, it is possible print them:

- `-view-dom, -dot-dom`
- `-view-postdom, -dot-postdom`

# Loop Information

Loop information are represented using two classes:

- `llvm::LoopInfo` analysis detects natural loops
- `llvm::Loop` represents a single loop

Using `llvm::LoopInfo` it is possible:

- navigate through top-level loops:  
`llvm::LoopInfo::begin(), llvm::LoopInfo::end()`
- get the loop for a given basic block:  
`llvm::LoopInfo::operator[] (llvm::BasicBlock *)`

# Loop Information


## Nesting Tree

Loops are represented in a **nesting tree**:

### Source

```
while(i < 10) {           // loop 1
  while(j < 10)           // loop 2
    while(k < 10)         // loop 3
    ...
  while(h < 10)           // loop 4
  ...
}
```

### Loop Nest



Nest navigation:

- children loops: `llvm::Loop::begin()`, `llvm::Loop::end()`
- parent loop: `llvm::Loop::getParentLoop()`

# Loop Information

## Query Loops

Accessors for relevant nodes also available:

```
pre-header llvm::Loop::getLoopPreheader()
header llvm::Loop::getHeader()
latch llvm::Loop::getLoopLatch()
exiting llvm::Loop::getExitingBlock(),
        llvm::Loop::getExitingBlocks(...)
exit llvm::Loop::getExitBlock()
     llvm::Loop::getExitBlocks(...)
```

Loop basic blocks accessible via:

```
iterators llvm::Loop::block_begin(),
          llvm::Loop::block_end()

vector std::vector<llvm::BasicBlock *> &llvm::Loop::getBlocks()
```



# Scalar Evolution

The **SC**alar **EV**olution framework:

- represents scalar expressions
- supports recursive updates
- lower burden of explicitly handling expressions composition
- is designed to support **general induction variables**

## Example

```
for.cond:
    %i.0 = phi [ 0, %entry ], [ %i.inc, %for.inc ]
    %cond = icmp ne %i.0, 10
    br %cond, label %for.body, label %for.end
for.inc:
    %i.inc = add nsw %i.0, 1
    br label %for.cond
for.end:
    ...
```

SCEV for %i.0:

- initial value 0
- incremented
- by 1 at each iteration
- final value 10

# Scalar Evolution

## Example

### Source

```
void foo() {
    int bar[10][20];

    for(int i = 0; i < 10; ++i)
        for(int j = 0; j < 20; ++j)
            bar[i][j] = 0;
}
```

SCEV {A,B,C}<D>:

- A initial
- B operator
- C operand
- D defining BB

### Induction Variables

```
%i.0 = phi i32 [ 0, %entry ], [ %inc6, %for.inc5 ]
--> {0,+,1}<nuw><nsw><%for.cond> Exits: 10
%j.0 = phi i32 [ 0, %for.body ], [ %inc, %for.inc ]
--> {0,+,1}<nuw><nsw><%for.cond1> Exits: 20
```

# Scalar Evolution

## More than Induction Variables

The scalar evolution framework manages **any scalar expression**:

### Pointer SCEVs

```
%arrayidx = getelementptr {...} %bar, i32 0, i32 %i.0
-->  {%bar,+,80}<nsw><%for.cond>
Exits: {%bar,+,80}<nsw><%for.cond>

%arrayidx4 = getelementptr {...} %arrayidx, i32 0, i32 %j.0
-->  {%bar,+,80}<nsw><%for.cond>,+,4}<nsw><%for.cond1>
Exits: {(80 + %bar),+,80}<nsw><%for.cond>
```

SCEV is an analysis used for common optimizations:

- induction variable substitution
- strength reduction
- vectorization
- ...

# Scalar Evolution

## SCEVs Design

SCEVs are modeled by the `llvm::SCEV` class:

- a subclass for each kind of SCEV: e.g. `llvm::SCEVAddExpr`
- instantiation disabled

A SCEV actually is a tree of SCEVs:

- $\{(80 + \%bar), +, 80\} = \{\%1, +, 80\}, \%1 = 80 + \%bar$

Tree leaves:

**constant** `llvm::SCEVConstant`: e.g. `80`

**unknown**<sup>9</sup> `llvm::SCEVUnknown`: e.g. `%bar`

SCEV tree explorable through the visitor pattern:

- `llvm::SCEVVisitor`

---

<sup>9</sup>Not further splittable

# Scalar Evolution

## Analysis Interface

The `llvm::ScalarEvolution` class:

- analyzes SCEVs for a `llvm::Function`

- builds SCEVs for values:

```
llvm::ScalarEvolution::getSCEV(llvm::Value *)
```

- creates new SCEVs:

```
llvm::ScalarEvolution::getConstant(llvm::ConstantInt *)
```

```
llvm::ScalarEvolution::getAddExpr(llvm::SCEV *, llvm::SCEV *)
```

```
...
```

- gets important SCEVs:

```
llvm::ScalarEvolution::getBackedgeTakenCount(llvm::Loop *)
```

```
llvm::ScalarEvolution::getPointerBase(llvm::SCEV *)
```

```
...
```

# Alias Analysis

Let  $X$  be an instruction accessing a memory location:

- is there another instruction accessing the same location?

Alias analysis tries to answer the question:

application memory operation scheduling  
problem often fails

Different algorithms for alias analysis:

- common interface – `llvm::AliasAnalysis` – for all algorithms
- by default, basic alias analyzer – `basicaa` – is used

## Requiring Alias Analysis

```
AU.addRequiredTransitive<llvm::AliasAnalysis>();
```

# Alias Analysis

## Memory Representation


### Source

```
%1 = load i16, i16* %a
%2 = load i16, i16* %b
store i16 %2, i32* %a
store i16 %1, i32* %b
```

### Distinct Locations



### Same Location



### Overlapping Locations



Basic building block is `llvm::AliasAnalysis::Location`:

- address: e.g. `%a`
- size: e.g. 2 bytes

# Alias Analyzer

## Basic Interface

Given two locations  $X$ ,  $Y$ , the alias analyzer classifies them:

- `llvm::AliasAnalyzer::NoAlias`:  $X$  and  $Y$  **are different** memory locations
- `llvm::AliasAnalyzer::MustAlias`:  $X$  and  $Y$  **are equal** – i.e. they points to the same address
- `llvm::AliasAnalyzer::PartialAlias`:  $X$  and  $Y$  **partially overlap** – i.e. they points to different addresses, but the pointed memory areas partially overlap
- `llvm::AliasAnalyzer::MayAlias`: **unable to compute** aliasing information – i.e.  $X$  and  $Y$  can be different locations, or  $X$  can be a complete/partial alias of  $Y$

Queries performed using:

- `llvm::AliasAnalyzer::alias(X, Y)`



# Alias Analyzer

## Mid-level Interface

Basic alias analyzer interface is low-level – we would like expressing queries about a single pointer  $X$ :

- how referenced memory location is accessed?
- which other instructions reference the same location?

What we need is a set, to classify memory locations:

- construct a `llvm::AliasSetTracker` starting from a `llvm::AliasAnalyzer *`
- it builds (one or more) `llvm::AliasSet`

For a given location  $X$ , a `llvm::AliasSet`:

- contains all locations aliasing with  $X$

# Alias Analyzer

## Alias Set Memory Accesses

Each alias set **references** the memory:

- `llvm::AliasSet::NoModRef`: no memory reference – i.e. the set is empty
- `llvm::AliasSet::Mod`: memory accessed in write-mode – e.g. a **store** is inside the set
- `llvm::AliasSet::Ref`: memory accessed in read-mode – e.g. a **load** is inside the set
- `llvm::AliasSet::ModRef`: memory accessed in read-write mode – e.g. a **load** and a **store** inside the set

# Alias Analyzer

## Mid-level Interface

Entry point is `llvm::AliasSetTracker::getAliasSetForPointer(...)`:

- `llvm::Value *`: location address
- `uint64_t`: location size
- `llvm::MDNode *`: used for type-based alias analysis <sup>10</sup>
- `bool *`: whether a new `llvm::AliasSet` has been created to hold the location – location does not alias up to now

Having the `llvm::AliasSet`:

- STL container-like interface: `size()`, `begin()`, `end()`, ...
- check reference type: `llvm::AliasSet::isRef()`, ...
- check aliasing type: `llvm::AliasSet::isMustAlias()`, ...

---

<sup>10</sup> set to NULL

# Memory Dependence Analysis

## Alias Analyzer High-level Interface

The `llvm::MemoryDependenceAnalysis` wraps alias analysis to answer queries in the following form:

- let `%foo` be an instruction accessing memory. Which preceding instructions does `%foo` depends on?

Reads:

- **store** instructions writing memory locations aliases with the one references by `%foo`

Writes:

- **load** instructions reading memory locations aliased with the one referenced by `%foo`

# Memory Dependence Analysis

## APIs

Let `%foo` be a `llvm::Instruction` accessing memory:

- call `llvm::MemoryDependenceAnalysis::getDependency(...)`
- you get a `llvm::MemDepResult`

Dependencies are classified:

- `llvm::MemDepResult::isClobber()`: an instruction clobbering – i.e. potentially modifying – location referenced by `%foo` has been found
- `llvm::MemDepResult::isDef()`: an instruction defining – e.g. writing – the exact location referenced by `%foo` has been found
- `llvm::MemDepResult::isNonLocal()`: no dependency found on `%foo` basic block
- `llvm::MemDepResult::isNonFuncLocal()`: no dependency found on `%foo` function

# Contents

# Conclusions

Inside LLVM there a lot of passes:

**normalization** put program into a canonical form

**analysis** get info about program

Please remember that

- a good compiler writer **re-uses** code
- check LLVM sources before re-implementing a pass

# Contents



# Bibliography I